

**CSCE 5300 Introduction to Big data and Data Science**

**ICE-8**

## Lesson Title: Spark

## Lesson Description: Spark with RDDs (transformation and actions) and Spark with data frames and SQL

## Lesson Overview:

Spark is a multi-language engine for executing data engineering, data science, and machine learning on single-node machines or clusters.

In Class Exercise:

1. Create spark RDD from external dataset(word\_list.txt)). Execute transformation and actions by scala.

Create RDD from external dataset(word\_list.txt).

Scala:

```
scala> val file = sc.textFile("/home/ljc/Downloads/word_list.txt")
file: org.apache.spark.rdd.RDD[String] = /home/ljc/Downloads/word_list.txt MapPartitionsRDD[1] at textFile at <console>:23

scala>
```

Python:

```
>>> file = spark.sparkContext.textFile("/home/ljc/Downloads/word_list.txt")
>>>
```

Change all words to uppercase and show the first two lines.

Scala:

```
scala> val myfile_up = file.map(line=>line.toUpperCase())
myfile_up: org.apache.spark.rdd.RDD[String] = MapPartitionsRDD[2] at map at <console>:23

scala> myfile_up.take(2)
res0: Array[String] = Array(
  THE PROJECT GUTENBERG ETEXT OF MOBY DICK BY HENRY JAMES
  WARD, COPYRIGHT LAWS ARE CHANGING ALL OVER THE WORLD, BE SURE TO CHECK)
```

Python:

```
>>> myfile_up = file.map(lambda line: line.upper())
>>> myfile_up.take(2)
['THE PROJECT GUTENBERG ETEXT OF MOBY WORD II BY GRADY WARD', 'COPYRIGHT LAWS ARE CHANGING ALL OVER THE WORLD, BE SURE TO CHECK']
>>>
```

Count the number of lines.

Scala:

```
scala> file.count()
res1: Long = 260
```

Python:

```
>>> file.count()
260
>>>
```

Count the number of word "PROJECT".

Scala:

```
scala> myfile_up.flatMap(line=>line.split(" ")).filter(c=>c.contains("PROJECT")).count()
res1: Long = 32
scala>
```

Python:

```
>>> myfile_up.flatMap(lambda line: line.split(" ")).filter(lambda word: word.count("PROJECT")).count()
32
>>>
```

Count the words in the dataset.

Scala:

```
scala> myfile_up.flatMap(line=>line.split(" ")).map(word=>(word,1)).reduceByKey((a,b)=>a+b).collect()
res0: Array[(String, Int)] = Array((MANAGE,1), (VOUL,1), (NOTES,1), (SEARCH,1), (FREQUENT,2), (OUR,2), (MIDNIGHT,1), (DICTIONARY,1), (THE,1), (NEVADA,2), ((OPTIONAL),1), (EXCLUSIONS,1), (READABLE,1), (DONATIONS,1), (DISPLAYS,1), (PREFER,1), (INDEMNIFY,1), (TAX,2), (ARCHIVE,3), (PLEASE,5), (DELTA,1), (ALLOW,1), (USERS,1), (NOVEL,1), (DEFECT,1), (LOOK,1), (GUTENBERG,1), (VALUE,1), (DESTINATION,1), (ELECTRONIC,1), (OPERATING,1), (CONSEQUENTIAL,1), (CHANGING,1), (UTILITY,1), ("PROJECT,3), (IS,18), (DATA,1), (REQUIREMENTS,2), (DAMAGED,1), (COMPRESSED,1), (REVISED,1), (FITNESS,1), (PEOPLE,2), ('ETUDE',1), (TRYING,1), (OWN,1), ((PLACES.TXT),1), (HOUR,2), (THOSE,3), (DECEMBER,1), (LARGE,1), (FOLLOWS,,1), (INTERESTED,1), (113,809,2), (CONSIDERABLE,1), (US,3), (IDEN...
```

Python:

```
>>> myfile_up.flatMap(lambda line: line.split(" ")).map(lambda word: (word,1)).reduceByKey(lambda a,b:a+b).collect()
res0: Array[(String, Int)] = Array((MANAGE,1), (VOUL,1), (NOTES,1), (SEARCH,1), (FREQUENT,2), (OUR,2), (MIDNIGHT,1), (DICTIONARY,1), (THE,1), (NEVADA,2), ((OPTIONAL),1), (EXCLUSIONS,1), (READABLE,1), (DONATIONS,1), (DISPLAYS,1), (PREFER,1), (INDEMNIFY,1), (TAX,2), (ARCHIVE,3), (PLEASE,5), (DELTA,1), (ALLOW,1), (USERS,1), (NOVEL,1), (DEFECT,1), (LOOK,1), (GUTENBERG,1), (VALUE,1), (DESTINATION,1), (ELECTRONIC,1), (OPERATING,1), (CONSEQUENTIAL,1), (CHANGING,1), (UTILITY,1), ("PROJECT,3), (IS,18), (DATA,1), (REQUIREMENTS,2), (DAMAGED,1), (COMPRESSED,1), (REVISED,1), (FITNESS,1), (PEOPLE,2), ('ETUDE',1), (TRYING,1), (OWN,1), ((PLACES.TXT),1), (HOUR,2), (THOSE,3), (DECEMBER,1), (LARGE,1), (FOLLOWS,,1), (INTERESTED,1), (113,809,2), (CONSIDERABLE,1), (US,3), (IDEN...
```

2. Create spark RDD from external dataset(shakespeare.txt). Execute transformation and actions by scala.

Change all words to lowercase and show the first 5 lines.

Count the total number of words.

Count the number of word "is".

Count the number of unique words in the dataset.

3. Create Spark dataframe from hotel\_booking data and execute some query.

Load data from the hotel\_booking.csv.

Scala:

```
scala> val df = spark.read.format("com.databricks.spark.csv").option("mode","DROP  
MALFORMED").option("inferSchema","true").option("header","true").csv("/home/ljc/  
Downloads/hotel_bookings.csv")  
df: org.apache.spark.sql.DataFrame = [hotel: string, is_canceled: int ... 30 mor  
e fields]
```

Python:

```
>>> df = spark.read.format("com.databricks.spark.csv").option("mode","DROPMALFORMED").option("header",True).option("inferSchema",True).csv("/home/ljc/Downloads/hotel_bookings.csv")  
>>>
```

Show some statistical values(mean, max value) of adults column.

Scala:

```
scala> df.describe("adults").show()  
+-----+-----+  
|summary|      adults|  
+-----+-----+  
|  count|      119390|  
|   mean|1.8564033838679956|  
| stddev|0.5792609988327523|  
|    min|              0|  
|    max|              6|  
+-----+-----+
```

Python:

```
>>> df.describe("adults").show()  
+-----+-----+  
|summary|      adults|  
+-----+-----+  
|  count|      119390|  
|   mean|1.8564033838679956|  
| stddev|0.5792609988327523|  
|    min|              0|  
|    max|              6|  
+-----+-----+
```

Count total number of canceled by hotel.

Scala:

```
scala> df.groupBy("hotel").sum("is_canceled").show()
+-----+-----+
|      hotel|sum(is_canceled)|
+-----+-----+
|  City Hotel|           33102|
|Resort Hotel|           11122|
+-----+-----+
```

Python:

```
>>> df.groupBy("hotel").sum("is_canceled").show()
+-----+-----+
|      hotel|sum(is_canceled)|
+-----+-----+
|  City Hotel|           33102|
|Resort Hotel|           11122|
+-----+-----+
```

Register the DataFrame as a global temporary view.

Scala:

```
scala> df.createGlobalTempView("hotelbook")
```

Python:

```
>>> df.createGlobalTempView("hotelbook")
```

Use query to count number of records is reservation\_status="canceled".

Scala:

```
scala> spark.sql("select count(reservation_status) from global_temp.hotelbook where reservation_status=\"Canceled\"").show()
+-----+
|count(reservation_status)|
+-----+
|              43017|
+-----+
```

Python:

```
>>> spark.sql("select count(reservation_status) from global_temp.hotelbook where reservation_status=\"Canceled\"").show()
+-----+
|count(reservation_status)|
+-----+
|              43017|
+-----+
```

Use query to count the number of agent group by hotel.

Scala:

```
scala> spark.sql("select sum(agent) from global_temp.hotelbook group by hotel").show()
+-----+
|sum(agent)|
+-----+
|2003876.0|
|6929877.0|
+-----+
```

Python:

```
>>> spark.sql("select sum(agent) from global_temp.hotelbook group by hotel").show()
+-----+
|sum(agent)|
+-----+
| 2003876.0|
| 6929877.0|
+-----+
```

Use query to count the number of babies when babies are greater than 0 by year.

Scala:

```
scala> val aa = spark.sql("select arrival_date_year, count(babies) from global_temp.hotelbook where babies>0 group by arrival_date_year").show()
+-----+-----+
|arrival_date_year|count(babies)|
+-----+-----+
|          2015|          213|
|          2016|          446|
|          2017|          258|
+-----+-----+
aa: Unit = ()
```

Python:

```
>>> spark.sql("select arrival_date_year, count(babies) from global_temp.hotelbook where babies>0 group by arrival_date_year").show()
+-----+-----+
|arrival_date_year|count(babies)|
+-----+-----+
|          2015|          213|
|          2016|          446|
|          2017|          258|
+-----+-----+
```

Use query to sort the number of canceled by country in decreasing order.

Scala:

```
scala> val aa = spark.sql("select country, sum(is_canceled) as all_canceled from global_temp.hotelbook group by country order by all_canceled desc").show()
+-----+-----+
|country|all_canceled|
+-----+-----+
|PRT| 27519|
|GBR| 2453|
|ESP| 2177|
|FRA| 1934|
|ITA| 1333|
|DEU| 1218|
|IRL| 832|
|BRA| 830|
|USA| 501|
|BEL| 474|
|CHN| 462|
|CHE| 428|
|NLD| 387|
|CNI| 254|
|RUS| 239|
|AUT| 230|
|SWE| 227|
+-----+-----+
```

Python:

```
>>> spark.sql("select country, sum(is_canceled) as all_canceled from global_temp.hotelbook group by country order by all_canceled desc").show()
+-----+-----+
|country|all_canceled|
+-----+-----+
|PRT| 27519|
|GBR| 2453|
|ESP| 2177|
|FRA| 1934|
|ITA| 1333|
|DEU| 1218|
|IRL| 832|
|BRA| 830|
|USA| 501|
|BEL| 474|
|CHN| 462|
|CHE| 428|
|NLD| 387|
|CNI| 254|
|RUS| 239|
|AUT| 230|
|SWE| 227|
+-----+-----+
```

## ICE Submission Guidelines

1. ICE Submission is individual.
2. ICE code has to be properly commented.
3. The documentation should include the screenshots of your code/queries and results.
4. Provide the explanation of the exercise for each question as per your understanding.
5. The similarity score for your document should be less than 15%.
6. Submit the source code (if any) properly commented and documentation (.pdf/.doc) with explanation and screenshot of source code/queries having input logic and output results.
7. Submission after the deadline is considered as late submission.

References:

<https://spark.apache.org/docs/latest/rdd-programming-guide.html#rdd-operations>

<https://spark.apache.org/docs/2.2.0/sql-programming-guide.html>

<https://sparkbyexamples.com/pyspark-rdd>

<https://sparkbyexamples.com/pyspark/different-ways-to-create-dataframe-in-pyspark/>